

MNIST DIGIT CLASSIFICATION THROUGH GPU





Prepared By:

- Muhammad Abdul Wahab Kiyani – 22i-1178 – CS-A
- Syed Ahmed Ali – 22i-1237 – CS-A

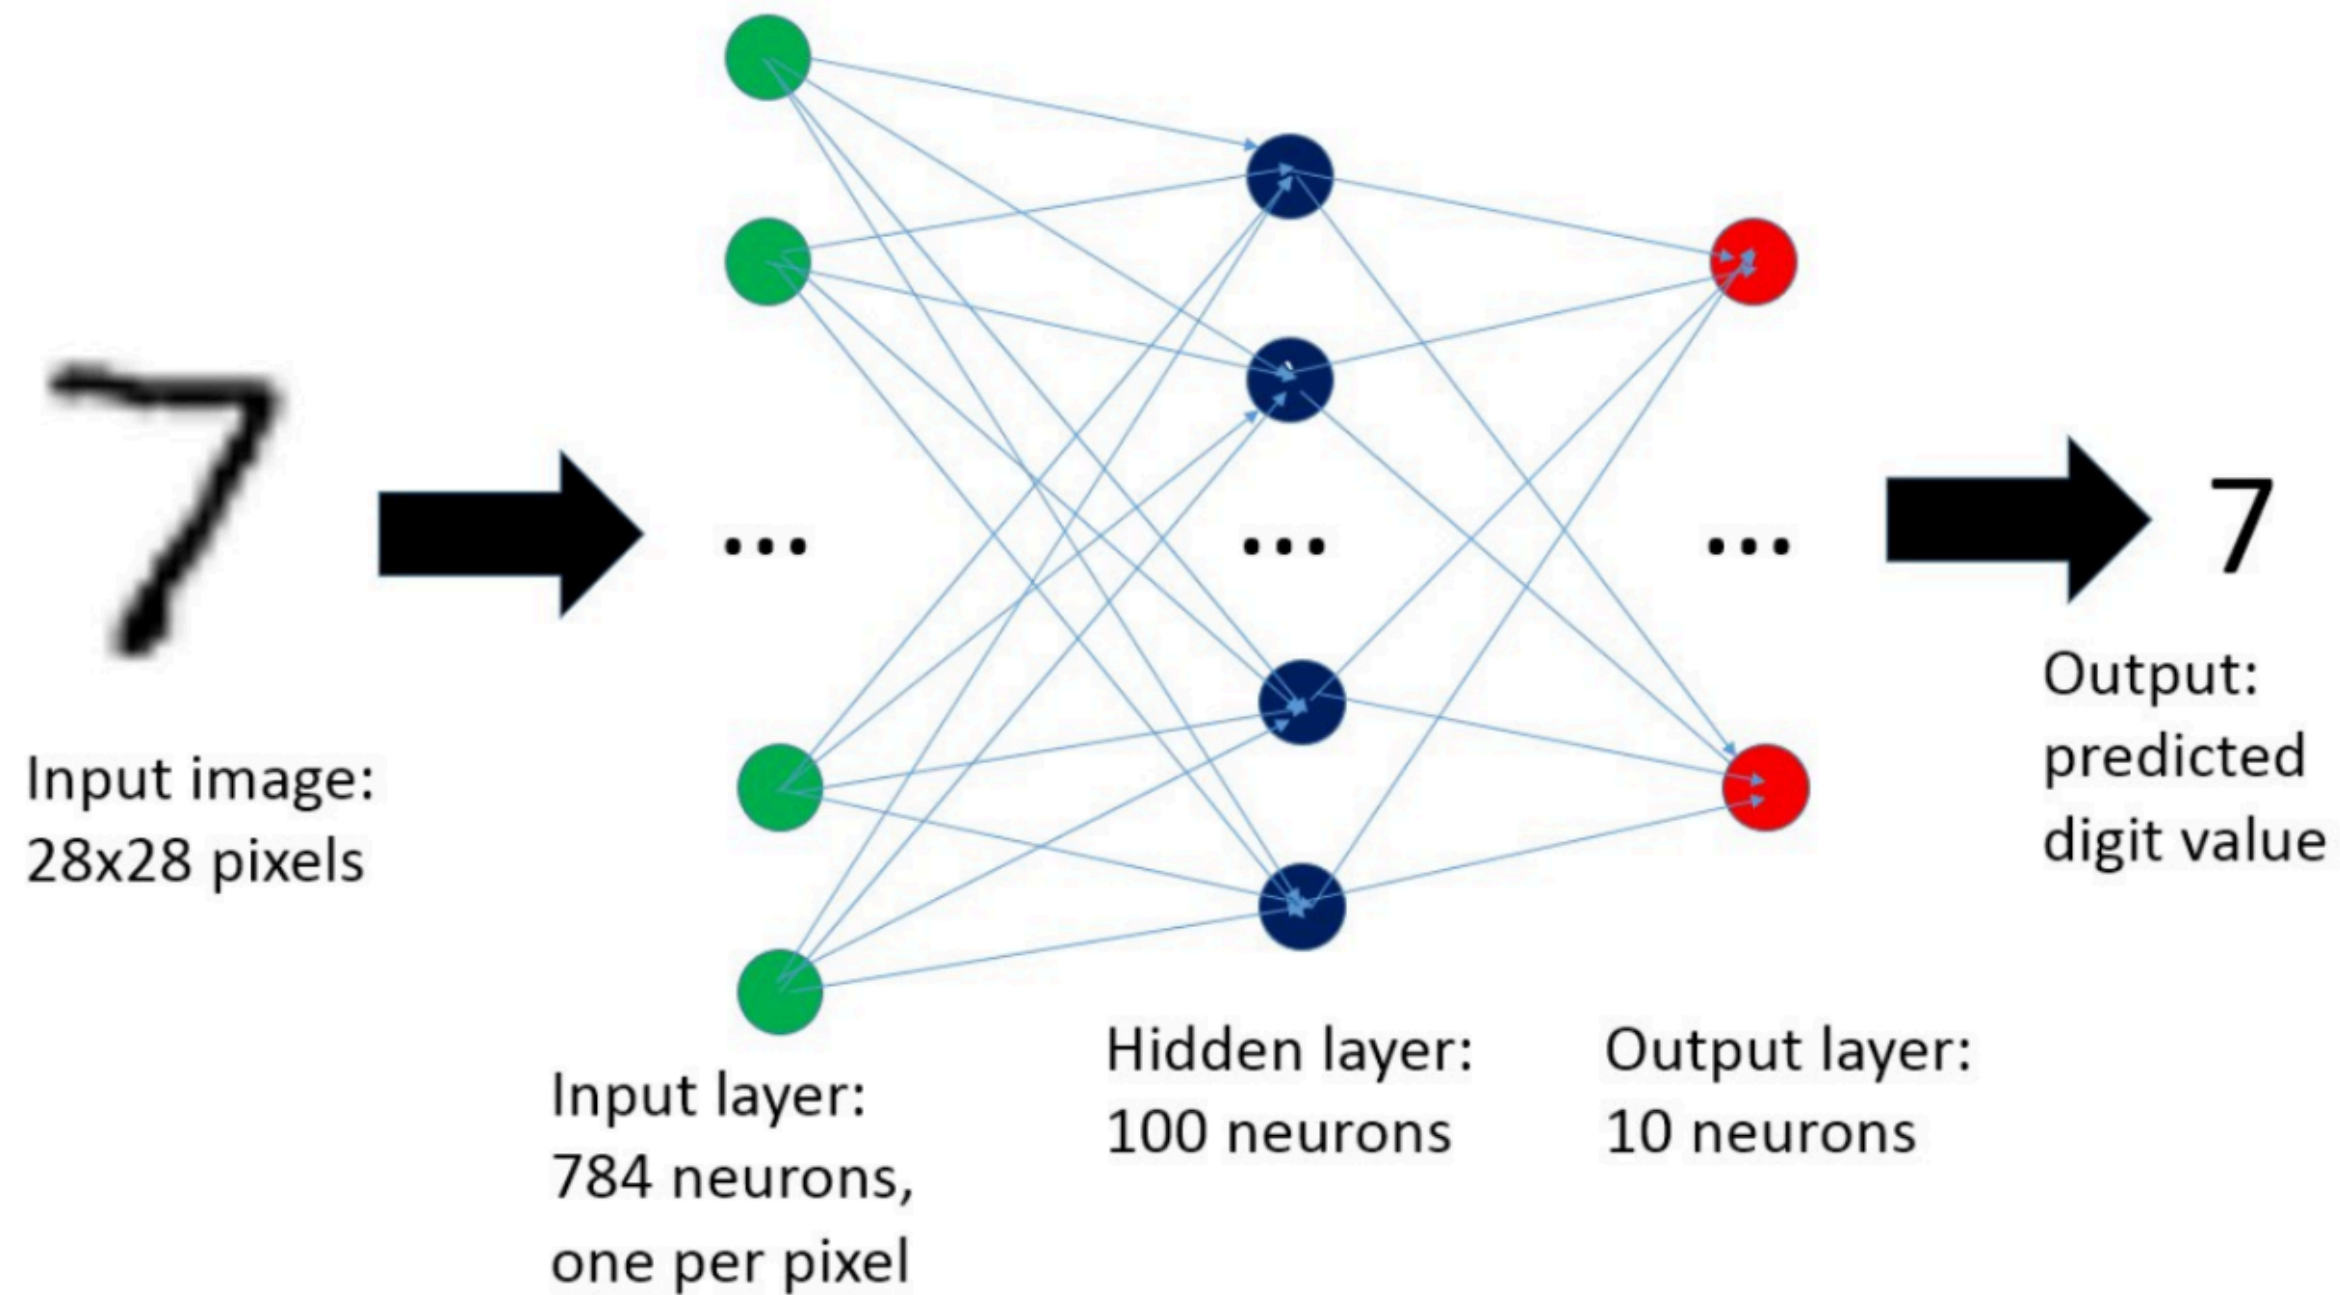
WHAT IS MNIST?

It's a dataset of handwritten digits (0–9) used to train and test machine learning models.

Input Data:






-  60,000 training images
-  10,000 test images
-  Each image: 28×28 grayscale pixels
-  Normalized to $[0, 1]$

NEURAL NETWORK



TESTING CONDITIONS



Model Parameters

-  Input Size: 784 (28×28 pixels)
-  Hidden Layer: 128 neurons
-  Output Size: 10
-  Epochs: 3
-  Learning Rate: 0.01

Tools Used

- Analysis: NVIDIA Nsight Systems
- Compilers: CUDA 12.8, GCC 13

Hardware Setup GPU

-  Development : RTX 3050 Ti
-  Benchmark : RTX 3080

Focus

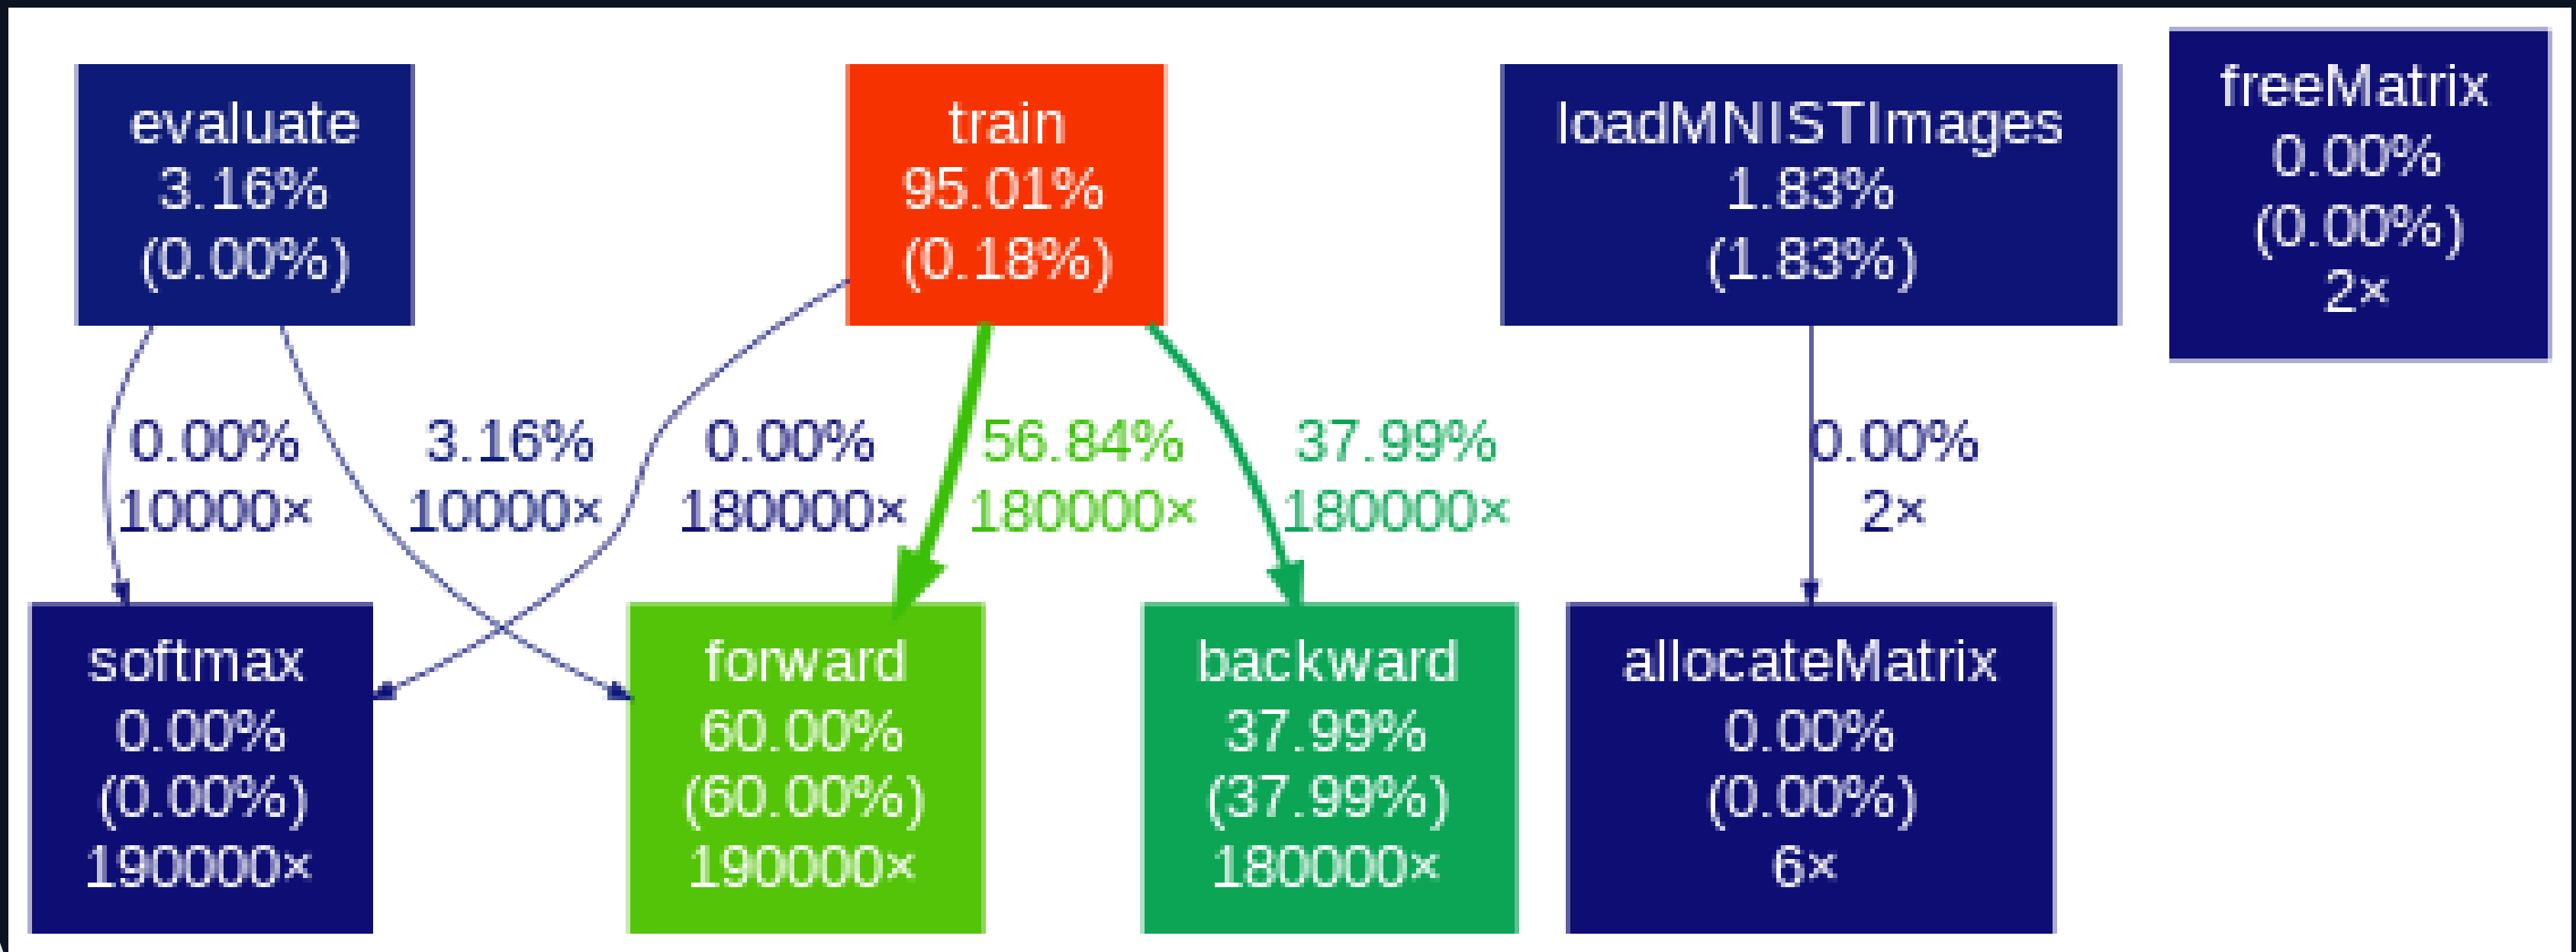
- Accuracy was the top priority
- All versions were trained for 3 epochs to ensure fair comparison

V1 — CPU BASELINE

- **Approach:** Classic sequential C++ loops
- **Performance:** 🐢 22.50s
- **Purpose:** Baseline to compare other versions

Simplest of all, but no parallelism

V1 — GPROF ANALYSIS



V2 — NAIVE CUDA

- **Approach:** Offload matrix ops to GPU
- **Issues:**
 - Uncoalesced memory access
 - Frequent host \leftrightarrow device transfers
 - No use of streams, shared memory, or smart load balancing

Performance: 🐢 41.46s

Code Snippet:

```
__global__ void update_weights_W2(double* d_W2, double* d_b2, double* d_hidden,
                                   double* d_d_output, double lr) {

    int i = blockIdx.x * blockDim.x + threadIdx.x;
    if (i < OUTPUT_SIZE) {
        for (int j = 0; j < HIDDEN_SIZE; j++) {
            int idx = i * HIDDEN_SIZE + j;
            d_W2[idx] -= lr * d_d_output[i] * d_hidden[j];
        }
        d_b2[i] -= lr * d_d_output[i];
    }
}
```


V3 — OPTIMIZED CUDA

- **Approach:**

- Tuned thread blocks
- Used shared memory
- Pinned host memory
- CUDA streams

Huge gains but complex.

Performance: ⚡ 3.32s

Code Snippet:

```
__global__ void update_weights_W2_shared(my_type* d_W2, my_type* d_b2, my_type* d_hidden,
                                         my_type* d_d_output, my_type lr) {

    __shared__ my_type shared_hidden[HIDDEN_SIZE];
    __shared__ my_type shared_d_output[OUTPUT_SIZE];
    int i = blockIdx.x * blockDim.x + threadIdx.x;

    for (int idx = threadIdx.x; idx < HIDDEN_SIZE; idx += blockDim.x) {
        shared_hidden[idx] = d_hidden[idx];
    }
    for (int idx = threadIdx.x; idx < OUTPUT_SIZE; idx += blockDim.x) {
        shared_d_output[idx] = d_d_output[idx];
    }
}
```

V4 — TENSOR CORES

- **Used:** FP16 + Tensor Core via WMMA
- **Issue:** Accuracy dropped from 96.85% to 70%
(but more epochs could've helped)

Observation:

- FP16 precision struggles with subtle weight updates during training.
- Rounding errors heavily impact backpropagation.

Code Snippet:

```
__global__ void tensor_matmul(const half_type* A, const half_type* B, my_type* C,
                             int M, int N, int K) {

    const int WMMA_M = 16;
    const int WMMA_N = 16;
    const int WMMA_K = 16;

    wmma::fragment<wmma::matrix_a, WMMA_M, WMMA_N, WMMA_K, half_type, wmma::row_major> a_frag;
    wmma::fragment<wmma::matrix_b, WMMA_M, WMMA_N, WMMA_K, half_type, wmma::col_major> b_frag;
    wmma::fragment<wmma::accumulator, WMMA_M, WMMA_N, WMMA_K, my_type> acc_frag;
```

V5 — OPENACC


- **Pros:**
 - Less manual effort
 - Best performance
- **Cons:**
 - Programmer has less control
 - Tricky debugging

Performance: ⚡ 2.24s

Code Snippet:

```
#pragma acc parallel loop gang vector
for (int i = 0; i < OUTPUT_SIZE; i++) {
    #pragma acc loop
    for (int j = 0; j < HIDDEN_SIZE; j++) {
        net->W2[i * HIDDEN_SIZE + j] -= LEARNING_RATE * d_output[i] * hidden[j];
    }
    net->b2[i] -= LEARNING_RATE * d_output[i];
}
```

SPEEDUPS



CONCLUSION

- GPU Parallelism drastically improves performance.
- Raw CUDA needs optimization for memory and communication efficiency.
- Memory management (shared memory, streams) is crucial for speedup.
- Tensor Cores offer great performance for matrix operations
- OpenACC simplifies parallelism with less manual effort but offers less control.